MSIN0097：Predictive Analytics
Individual Coursework

Candidate number: VRDK8

Word Count: 1988

1. Obtain a dataset and frame the predictive problem

1.1 Predictive Task and Target Definition

This project formulates a supervised binary classification task to predict whether a client subscribes to a term deposit (deposit). The prediction is defined under a strict pre-contact setting, meaning only information available before campaign outcome is used. This timing assumption is explicitly enforced to prevent target leakage and reflect realistic deployment conditions.

The target variable is encoded as:

- 1 = subscription ("yes")
- 0 = no subscription ("no")

The model outputs a probability $P(\text{deposit} = 1)$, which can later be converted into a classification decision using an operational threshold.

1.2 Success Metrics and Constraints

Because subscription outcomes are moderately imbalanced, overall accuracy is not used as the primary metric. The main evaluation metric is PR-AUC for the positive class, complemented by Recall, Precision, and Balanced Accuracy. Threshold-based metrics are used to reflect campaign trade-offs.

Fairness is evaluated across age and marital groups. Equal Opportunity (TPR disparity) is treated as the primary fairness metric, with Demographic Parity and Equalized Odds reported for diagnostic comparison. Operational constraints include interpretability, near real-time scoring feasibility, asymmetric error costs, and exclusion of post-outcome variables (e.g., call duration) to avoid leakage.

1.3 Assumptions and Limitations

The dataset reflects historical campaign behaviours and may encode prior operational bias. Fairness metrics identify statistical disparities but do not imply causality. Results depend on group definitions and threshold selection and may not generalize across time or institutions.

1.4 Agent tooling

Codex was used to sanity-check task framing (pre-contact prediction) and metric choices. All timing assumptions (feature availability) and evaluation criteria (PR-AUC, EO fairness) were finalized after manual verification against the dataset schema and deployment constraints.

2. Explore the data to gain insights

2.1 Dataset Overview and Target Distribution

• Figure 1. Class Distribution of Target (Deposit Subscription)

The dataset contains 11,162 observations and 17 variables. The target variable is deposit, indicating whether a client subscribed to a term deposit.

The class distribution is relatively balanced, with the majority-class baseline accuracy approximately 52.6%. Although imbalance is not severe, accuracy alone would not sufficiently reflect performance on the positive class. Therefore, imbalance-aware metrics (e.g., precision, recall, PR-AUC) remain appropriate for evaluation.

2.2 Data Quality and Implicit Missingness

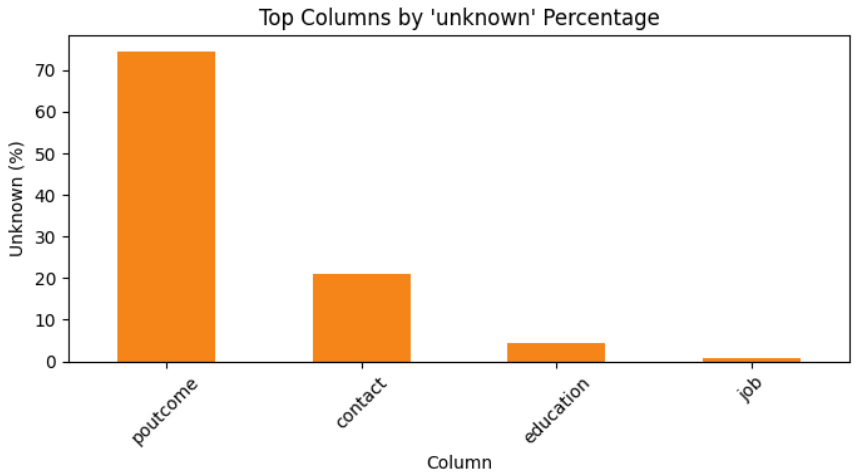

• Figure 2. Proportion of "Unknown" Category by Feature

No formal missing values (NaN) were detected. However, several categorical variables contain a substantial proportion of "unknown" placeholder values. In particular, poutcome and contact exhibit high rates of unknown entries.

This indicates that missingness is encoded as a category rather than as null values. Such placeholders may influence model learning and interpretation and should be treated carefully in downstream preprocessing.

2.3 Leakage Investigation (Pre-contact Setting)

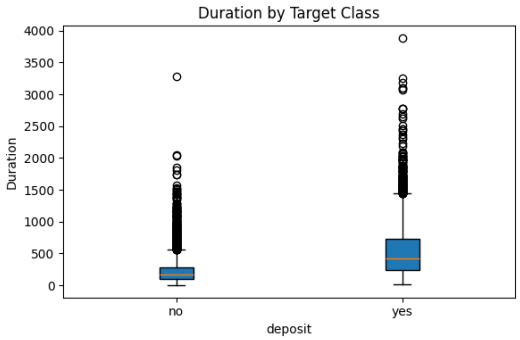

- Figure 3. Call Duration Distribution Across Target Classes

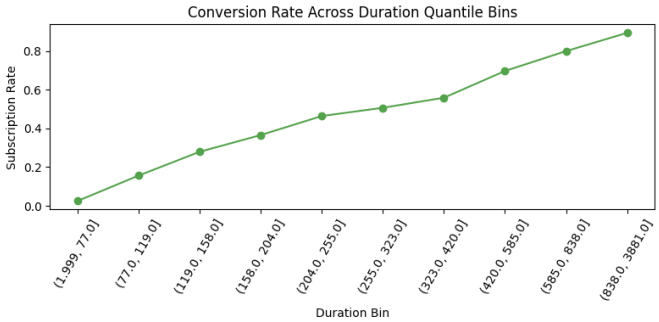

- Figure 4. Subscription Rate Across Call Duration Quantile Bins

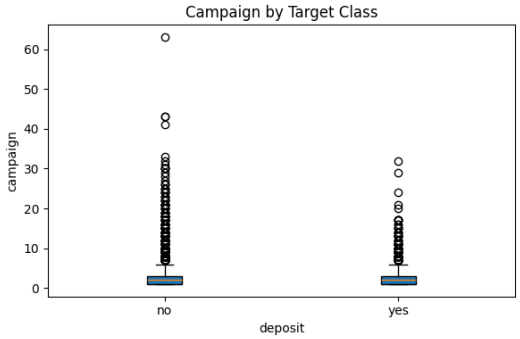

- Figure 5. Campaign Frequency Distribution Across Target Classes

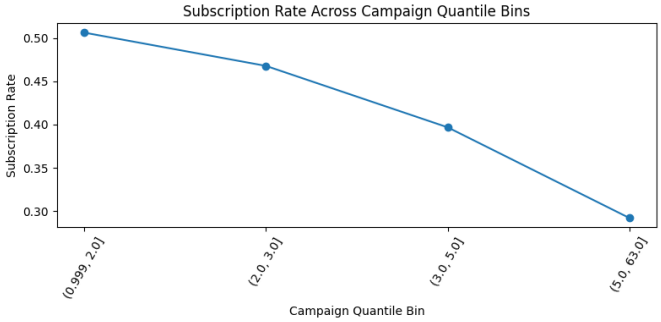

- Figure 6. Subscription Rate Across Campaign Quantile Bins

A strong relationship is observed between duration (call duration) and deposit. Clients who subscribed to the deposit tend to have significantly longer call durations. Moreover, subscription rates increase almost monotonically across duration quantile bins, with the highest bins approaching very high conversion rates. However, call duration is only known after the call has been completed. In a pre-contact prediction setting, including duration would introduce post-outcome information and artificially inflate model performance. Therefore, duration must be excluded from modelling to avoid data leakage.

Unlike duration, which clearly represents post-outcome information, campaign presents a potential timing ambiguity. The distribution and long-tail behaviour suggest it accumulates across contact attempts. While not inherently a leakage feature, its eligibility depends on whether it reflects only prior contacts at the prediction moment. This temporal dependency will need to be clarified before feature inclusion decisions are made in the modelling stage.

2.4 Numerical Distributions and Outliers

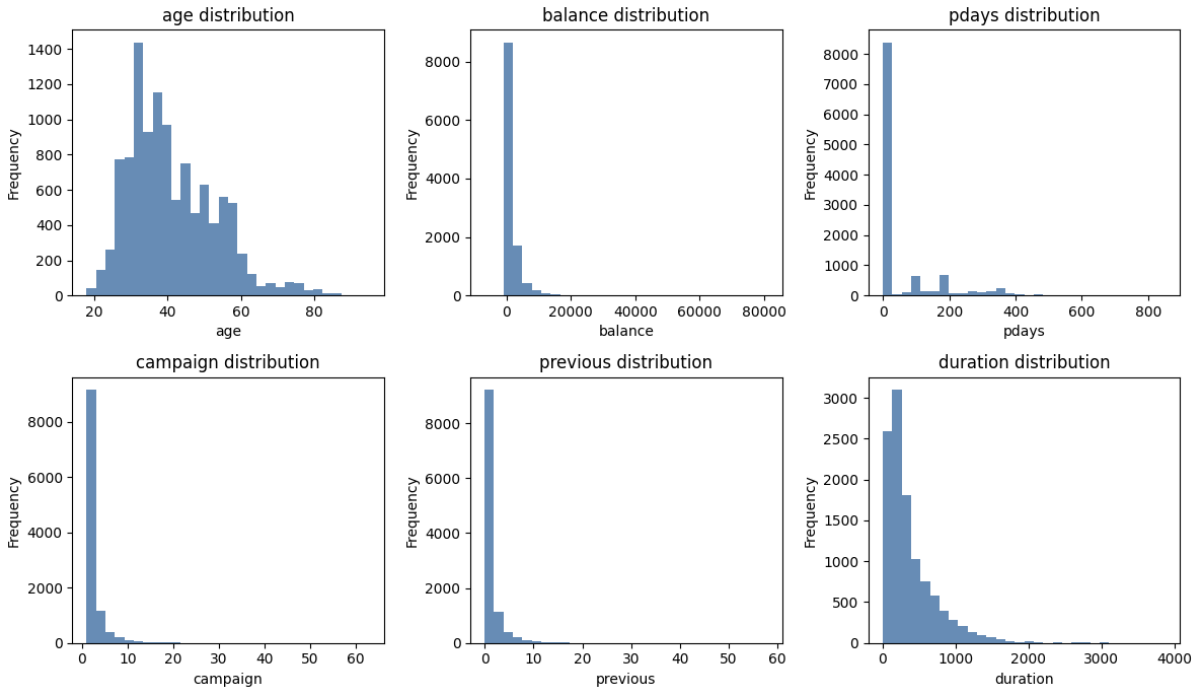

- Figure 7. Distributional Characteristics of Numerical Variables

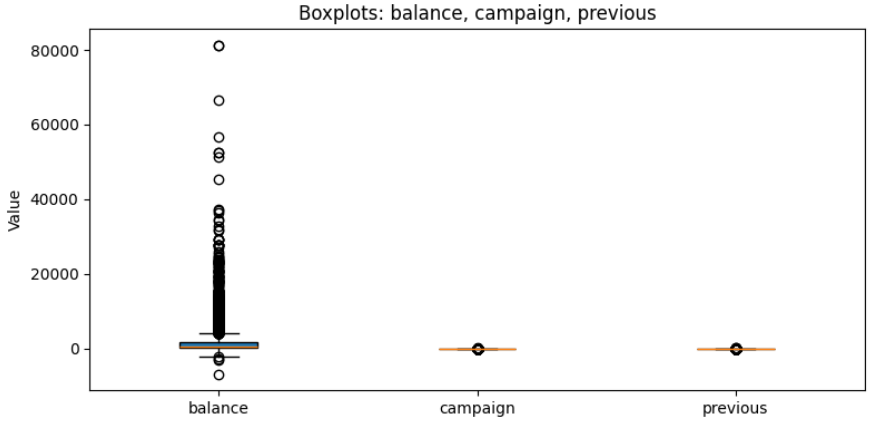

- Figure 8. Outlier Analysis for Balance, Campaign, and Previous Contacts

Several numerical variables (e.g., balance, campaign, previous) exhibit strong right skewness and long-tailed distributions. Extreme values are present but appear structurally plausible within the business context rather than indicative of data errors.

The variable pdays contains a sentinel value of -1, representing clients who were not previously contacted in the recent campaign. Approximately three-quarters of observations contain this value. Therefore, -1 should be interpreted as a distinct state rather than a standard continuous lag value.

These distributional characteristics suggest that model choice and feature handling must account for skewness and sentinel-coded behavior.

2.5 Categorical Structure

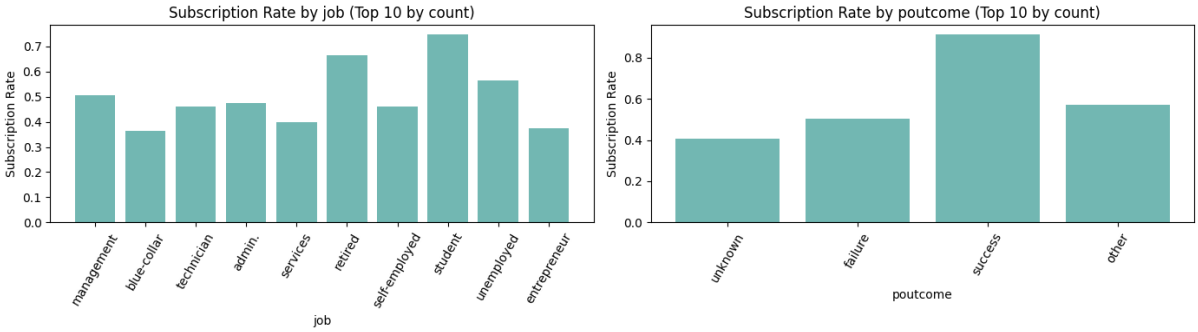

- Figure 9. Subscription Rate by Job and Previous Campaign Outcome

Categorical variables such as job and poutcome demonstrate meaningful variation in subscription rates. In particular, poutcome shows strong differentiation, with clients whose previous outcome was successful exhibiting substantially higher subscription rates.

Some categories contain relatively small sample sizes, which may result in unstable empirical rates. Rare-level effects should therefore be interpreted cautiously.

2.6 Fairness-Oriented Base Rate Analysis

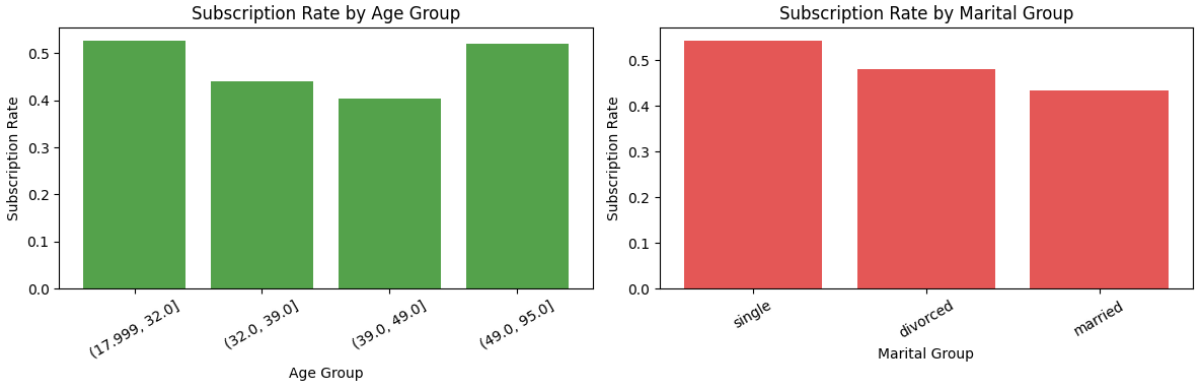

- Figure 10. Subscription Rate by Age and Marital Status

Base subscription rates differ across age groups and marital status categories. This indicates heterogeneity in underlying positive outcome rates between demographic groups.

Such differences have implications for fairness evaluation. When group-level base rates differ, fairness definitions such as Demographic Parity and Equal Opportunity may conflict. Therefore, fairness metric selection must reflect observed base-rate heterogeneity.

2.7 Subgroup Support

All age and marital groups contain sufficient sample sizes to support descriptive fairness analysis. However, subgroup-level conclusions should always be interpreted alongside support counts to avoid over-interpreting small-group effects.

2.8 Agent tooling

The agent suggested candidate leakage features (e.g., duration/campaign) and EDA checks; I validated them by inspecting target-rate patterns (e.g., duration bins) and by confirming whether each feature is observable at the pre-contact decision point.

3. Prepare the data

3.1 Prediction Setting and Feature Eligibility

This project adopts a strict pre-contact prediction setting, where subscription probability is estimated before the first contact attempt. Only variables available prior to outreach are considered eligible.

Consistent with the leakage findings in Step 2, duration is excluded as it represents post-call information. The variable campaign is also excluded under the first-contact definition, since it accumulates across attempts and may not be known at the initial scoring point.

Two feature sets are defined to reflect timing assumptions:

- **Set A (baseline strict):** excludes duration and campaign (14 predictors).
- **Set B (extra strict):** additionally excludes day, month, and contact (11 predictors), which may depend on campaign scheduling decisions.

This specification explicitly separates confirmed leakage (duration) from timing-sensitive variables, allowing sensitivity analysis under stricter operational assumptions.

3.2 Data Splitting and Experimental Protocol

The dataset (11,162 observations) is split into training, validation, and test sets using stratified sampling (approximately 64/16/20). Class proportions (≈52.6% "no", 47.4% "yes") are preserved across splits. Feature eligibility decisions are finalized before model training. The held-out test set is reserved strictly for final evaluation to prevent optimistic bias.

3.3 Feature Screening and Proxy Leakage Check

To detect potential proxy leakage, association diagnostics were computed for retained predictors, including mutual information ranking, numeric correlations, and categorical target-rate spread.

The most strongly associated variables (e.g., poutcome, month, balance, pdays, previous) reflect historical campaign signals or financial characteristics rather than post-outcome information. Month and poutcome show relatively large category-rate spreads (about 0.58 and 0.51), so they are treated as timing-sensitive risk signals and monitored in the strict-setting comparison, rather than being removed without validation.

These diagnostics were conducted prior to model training for exploratory leakage detection only. They did not influence hyperparameter tuning or model selection, and the test set remained untouched during this process.

3.4 Preprocessing Pipeline and Data Validation

The target variable is explicitly encoded as binary (1 = subscription, 0 = non-subscription) to ensure consistent metric computation.

Schema checks confirm that excluded variables (e.g., duration, campaign) and the encoded target are not present in the feature matrix. Stratified splits are applied prior to model fitting, and class distributions are verified to ensure consistency across subsets.

Data-quality diagnostics show no missing values across columns in the current dataset snapshot. Categorical unknown values are concentrated in poutcome (74.59%), contact

(21.02%), and education (4.45%), and these patterns were retained and handled in downstream preprocessing rather than dropped ad hoc.

3.5 Agent tooling

The agent assisted in drafting the scikit-learn split structure and preprocessing outline. I independently verified correctness by enforcing schema checks (ensuring excluded variables were not present in X), confirming stratified splits and class proportions, and explicitly encoding the target as 0/1 to ensure consistent positive-class handling in downstream metrics.

4. Explore different models and shortlist the best ones

4.1 Baseline Model

Logistic regression was adopted as the baseline model for this binary classification task. It provides a transparent and interpretable benchmark, consistent with the project's emphasis on leakage control and fairness analysis. Model comparison was conducted using PR-AUC for the positive class, as this metric better reflects precision–recall trade-offs than accuracy under mild class imbalance.

Hyperparameters for Logistic Regression were tuned using a small grid search over regularization strength and class weighting, evaluated via cross-validation.

4.2 Model Comparison and Selection

Model comparison was performed using 5-fold stratified cross-validation on the training split only, with PR-AUC (average precision) as the evaluation metric. No validation or test data were used during hyperparameter tuning.

Random Forest and HistGradientBoosting were tuned using randomized search to efficiently explore larger hyperparameter spaces. Among the compared models, HistGradientBoosting achieved the highest mean CV PR-AUC (0.794), outperforming Random Forest (0.784) and Logistic Regression (0.765), while maintaining low fold-to-fold variability. The final model was selected strictly based on mean CV PR-AUC. The held-out test set was reserved exclusively for final evaluation. On the test set, HistGradientBoosting achieved PR-AUC = 0.777, indicating a modest and expected generalization gap relative to cross-validation performance.

Hyperparameter ranges were initially proposed by the agent and subsequently validated via cross-validation experiments before final model selection.

4.3 Agent tooling

The agent suggested model families and compact hyperparameter ranges; I accepted or refined them based on 5-fold stratified cross-validation (PR-AUC) results, and kept the test set strictly untouched until after final model selection.

5. Fine-tune and evaluate

5.1 Tuning Strategy and Robust Evaluation

After model selection, a classification threshold was determined using the validation set by maximizing F1 score (threshold = 0.3368). This threshold was then applied once to the test set for final evaluation, ensuring that threshold tuning did not rely on test data.

During implementation, an agent-generated evaluation step produced a pos_label inconsistency error due to mismatched target encoding (string vs numeric labels). This was identified by checking label types and resolved by explicitly encoding the target as 0 (no) and 1 (yes), ensuring consistent positive-class handling across all metrics.

5.2 Test Performance (Single Use)

Threshold-independent metrics:

- PR-AUC: 0.7922
- ROC-AUC: 0.7915

Threshold-dependent metrics (validation-selected threshold):

Confusion matrix:

$$\begin{bmatrix} 654 & 521 \\ 186 & 872 \end{bmatrix}$$

- Precision: 0.6260
- Recall: 0.8242
- F1: 0.7115

The model favours recall under the chosen F1-optimized threshold, increasing true positive detection at the cost of higher false positives.

5.3 Calibration and Error Analysis

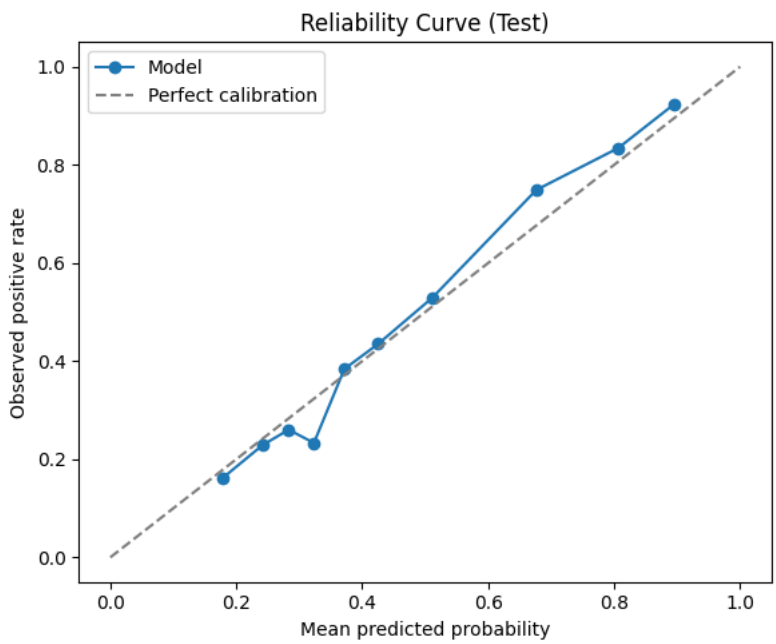

- Figure 11. Reliability Curve on the Test Set

Calibration was assessed via reliability curve and Brier score (0.1828), indicating reasonably aligned predicted probabilities with observed outcomes, with minor mid-range deviations.

Inspection of high-confidence errors revealed that false positives frequently exhibited prior campaign success (poutcome = success), suggesting strong reliance on historical interaction signals. False negatives often involved limited prior contact history (poutcome = unknown, pdays = -1), indicating weaker signals reduce predicted probability even when subscription occurs. The model appears to over-weight prior campaign success as a predictive signal, which may reflect historical correlation but not necessarily causal stability. This suggests potential sensitivity to distribution shift in future campaign cycles.

Overall, the evaluation framework supports the conclusion that performance is consistent with genuine generalization under a disciplined experimental protocol.

5.4 Agent tooling

I documented and corrected an agent-generated implementation error: a positive-class encoding mismatch that triggered a pos_label error. I identified it via y.unique() and resolved it by enforcing explicit binary encoding (0=no, 1=yes) before rerunning evaluation. Final test metrics were computed once after model selection

6. Present the final solution

Based on cross-validation comparison, HistGradientBoosting was selected as the final model, achieving the highest and most stable PR-AUC among the evaluated approaches. The generalization gap between cross-validation (≈0.792) and test performance (≈0.777–0.792 depending on metric) was modest, indicating no severe overfitting under the strict pre-contact setting. Calibration analysis further suggested that predicted probabilities are reasonably aligned with observed outcomes, supporting their use for ranking-based targeting decisions.

This model is intended to estimate the probability of term deposit subscription prior to the first campaign contact. It is suitable for pre-contact prioritization but not for real-time call optimization, automated rejection decisions, or causal interpretation of customer behavior. All features were filtered according to explicit prediction-time availability to prevent leakage, and evaluation followed strict train/validation/test separation.

Several limitations remain. The dataset reflects historical marketing practices and may encode prior operational biases. Performance depends on threshold selection and business objectives, and the chosen F1-optimized threshold emphasizes recall at the cost of higher false positives. Additionally, results are based on a single dataset and may not generalize across time without temporal validation. Future improvements could include time-based validation, cost-sensitive optimization, and formal fairness auditing.

7. Appendix

bank.csv

I am using bank.csv (11,162 rows, 17 columns) for my coursework.

I want to complete Step 1: Problem Framing.

Please help me structure this section clearly, but do not generate code.

1. Define the predictive task precisely, including the prediction timing (what is known at prediction time).
2. Define the target variable and prediction type.
3. Propose appropriate primary and secondary success metrics, considering likely class imbalance.
4. Define fairness metrics for age and marital status, and explain the trade-offs between them.
5. Define practical constraints (interpretability, latency, business cost of errors).
6. Identify high-level assumptions and limitations of this predictive setup.
7. Explicitly state what AI (Codex) can assist with and what must be independently verified by me.

Important:
- Do not assume all variables are usable at prediction time.
- Do not provide implementation details.
- Keep the tone academic and suitable for a 2000-word report.
- Avoid inventing dataset statistics.

- Appendix 1. Agent tooling for Step 1

Based on the bank dataset structure, what initial EDA checks should I perform to understand class imbalance, missing values, and categorical distribution before modeling?

Do not write code; just list diagnostic steps and explain why each matters.

1. **Confirm target encoding and class counts (`deposit`)**
   Why it matters: establishes the baseline imbalance level and prevents mistakes (for example, inconsistent labels like `yes/no`, whitespace variants, or case differences).
2. **Compute class proportion and majority-class baseline performance**
   Why it matters: shows how misleading raw accuracy could be and sets a naive benchmark any model must beat.
3. **Check class balance across key segments (`month` , `contact` , `job` , `marital`)**
   Why it matters: reveals whether imbalance is uniform or concentrated in subpopulations, which affects sampling strategy, evaluation, and fairness interpretation.

- Appendix 2. Agent tooling for Step 2 EDA checks selection

I am performing Step 2 (Exploratory Data Analysis) for bank.csv.

Goal:
Produce clean, reproducible Python EDA code (pandas + matplotlib only) organized into clearly separated modules.
Do NOT train any models or create preprocessing pipelines.

Use this structure:

Module 1: Dataset Overview & Target Imbalance

- Print shape, columns, dtypes.
- Detect whether target column is named 'deposit' or 'y'.
- Show class counts and proportions.
- Compute majority-class baseline accuracy.
- Create one bar chart of class distribution.
- Add short markdown explaining why imbalance matters for metric choice.

- Appendix 3. Agent tooling for Step 2 EDA module codes

I suspect that the campaign variable may introduce temporal leakage in a pre-contact prediction setting.
In this dataset, campaign represents the number of contacts performed during the current campaign for a given client.
Please help me perform a leakage-focused analysis of campaign similar to what we did for duration, with the following structure:
Descriptive statistics of campaign by target (deposit)
Boxplot of campaign by target class
Subscription rate across campaign quantile bins
Correlation between campaign and duration
Short narrative interpretation:
Is campaign likely known at prediction time?
Does it appear to accumulate during the campaign?
Could it encode post-contact information?
Use pandas and matplotlib only.
Keep it notebook-friendly (no argparse, no CLI arguments).

- Appendix 4. Agent tooling for Step 2 EDA further analysis

Before proceeding to modeling, please print the final list of feature columns used in X.
For each feature, briefly classify whether it is:
- demographic
- financial
- historical campaign
and confirm that it is available prior to first contact under a strict pre-contact setting.
Do not modify the dataset

python

```python
# %% [markdown]
# ### Feature audit before modeling (no dataset modification)
# Prints final feature list from `X`, classifies each feature, and confirms pre-contact avai

# %%
import pandas as pd
```

- Appendix 5. Agent tooling for Step 3 for Pre-modeling Feature Audit and Eligibility Verification

Check whether any remaining features show unusually high correlation with the target.
Compute mutual information or simple association metrics for all features. Flag any feature that appears suspiciously predictive.
Do not train a model yet.

python

```python
# %% [markdown]
# ### Leakage-oriented feature association scan (no model training)
# Goal:
# - Compute association strength between each feature and target (`y_bin`)
# - Use mutual information across all features
# - Add simple association metrics (numeric correlation, categorical rate spread)
# - Flag features that look suspiciously predictive
```

- Appendix 6. Agent tooling for Step 3 for Leakage-Oriented Association Scan

Please create two feature sets for strict pre-contact modeling:
Set A (baseline strict): drop duration and campaign (current approach).
Set B (extra strict): additionally drop day, month, and contact due to potential planning/time leakage ambiguity.

For each set, print the final feature list and X shape.
Do not train any model yet.

python

```python
# %% [markdown]
# ## Create Two Feature Sets for Strict Pre-Contact Modeling
# - Set A (baseline strict): drop `duration`, `campaign`
# - Set B (extra strict): drop `duration`, `campaign`, `day`, `month`, `contact`
# - Print final feature list and X shape for each
# - No model training
```

- Appendix 7. Agent tooling for Step 3 for building feature sets

I plan to compare Logistic Regression, Random Forest, and HistGradientBoosting on the bank marketing dataset for step 4 of the project
Suggest reasonable hyperparameter ranges for a small, computationally efficient grid search for each model.
Do not suggest extreme tuning; keep it practical for coursework.

For each model:
- Explain why those ranges make sense for this dataset size.
- Keep the grid small (3–4 values per key parameter).

For ~11k rows, a **small, structured grid** is enough to compare model families without overfitting your coursework budget.

Logistic Regression (baseline, fast, interpretable)
Use a compact grid around regularization strength and imbalance handling:

- Appendix 8. Agent tooling for Step 4 for Computationally Efficient Hyperparameter Grid Design

I will verify these hyperparameter suggestions with experiments.
Please generate VSCode Jupyter notebook code (no CLI) that does the following:
1) Uses my existing preprocessing pipeline (ColumnTransformer + OneHotEncoder + scaling) and the strict pre-contact Feature Set A as the default.
2) Runs model selection using cross-validation with:
  - StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
  - scoring = "average_precision" (PR-AUC)
3) For Logistic Regression: use a small GridSearchCV (C and class_weight only; keep solver fixed to avoid sparse issues).
4) For RandomForestClassifier and HistGradientBoostingClassifier: use RandomizedSearchCV with a fixed budget (e.g., n_iter=30 each) based on the suggested ranges.
5) Print a compact results table for each model:
  - best mean CV PR-AUC, std CV PR-AUC, best_params_
6) Create a final comparison table across the three models and shortlist the best model based on mean CV PR-AUC.

Important:
- Keep the test set untouched (use only training data for CV).
- Use the same random_state everywhere for reproducibility.

- Appendix 9. Agent tooling for Step 4 for Strategic Hyperparameter Search Space Design

I'm moving to Step 5 (tuning + robust evaluation + error analysis).

I already did Step 4 model selection with 5-fold Stratified CV using PR-AUC. Best model is HistGradientBoosting. I have a strict "test set untouched" policy: use the test set exactly once after selecting the best model.
My splits already exist in memory: X_train_A, y_train_A, X_val_A, y_val_A, X_test_A, y_test_A (all y are binary 0/1; 1 is the positive class).

Task:
1) Fit the selected best pipeline on (X_train_A + X_val_A) and evaluate once on X_test_A:
  - Report TEST PR-AUC (average precision).
  - Also report ROC-AUC.
2) Choose a decision threshold:
  - Either 0.5 or pick the threshold that maximizes F1 on the validation set (state which you did).
  - Then compute confusion matrix + precision, recall, F1 on the test set.
3) Add calibration diagnostics on the test set:
  - Reliability curve (calibration plot)
  - Brier score
4) Error analysis / failure modes:
  - Extract top 20 false positives and top 20 false negatives by predicted probability.
  - For each group, summarize key feature distributions (at least: poutcome, pdays, previous, balance).
  - Give 2–3 short qualitative observations about patterns.

Please return clean, runnable sklearn code cells (no re-loading CSV, no re-splitting). Assume the split artifacts already exist.

- Appendix 10. Agent tooling for Step 4 for Final Model Evaluation, Calibration, and Error Analysis

#	Step 1 Element	Agent Suggestion	Decision	Verification	Rationale
1	Predictive task & timing	Binary classification predicting deposit; define pre-contact decision point	Accept	Reviewed variable descriptions to confirm prediction-time assumptions	Explicit timing prevents target leakage
2	Feature eligibility	Exclude post-outcome variables (e.g., duration)	Accept	Confirmed duration known only after call completion	Including it would inflate offline performance unrealistically
3	Target definition	deposit (yes/no) → binary classification	Accept	Checked dataset schema and label encoding	Matches supervised learning setup
4	Performance metrics	PR-AUC, Recall, Precision, F1, ROC-AUC	Modify	Selected PR-AUC as primary; limited secondary metrics	Avoid metric redundancy; focus on imbalance-aware evaluation
5	Fairness metrics	Demographic Parity, Equal Opportunity, Equalized Odds	Modify	Chose Equal Opportunity as primary fairness metric	TPR parity aligns better with marketing targeting context
6	Group definition	Create age and marital subgroups	Modify	Will justify grouping based on data distribution	Avoid arbitrary demographic thresholds
7	Practical constraints	Interpretability, latency, asymmetric error costs	Accept	Incorporated into evaluation framing	Ensures deployment realism
8	AI tooling plan	Codex assists; student verifies	Accept	Documented independent validation steps	Demonstrates supervision and methodological ownership
#	**Step 2 Element**	**Agent Suggestion**	**Decision**	**Verification**	**Rationale**
1	Target balance	Compute baseline accuracy	Accept	Reviewed proportions	Supports PR-based evaluation
2	Missing data	Check NaN only	Modify	Found "unknown"placeholders	Prevent hidden missingness bias
3	Leakage	Use duration as predictor	Reject	Verified post-call timing	Avoid target leakage
4	pdays encoding	Treat as numeric	Modify	High frequency of -1	Sentinel-coded state
5	Outliers	Flag via IQR	Accept	Observed skewness	Inform model choice

6	Demographic rates	Compare age/marital	Accept	Base rates differ	Fairness metric implications
7	Timing ambiguity check (campaign)	Run leakage-focused diagnostics for campaign (by target stats, boxplot, quantile-bin subscription rate, corr with duration)	Accept	Confirmed: subscription rate decreases as campaign bin increases; campaign has long-tail; correlation with duration near 0	campaign is not clear post-outcome like duration, but pattern suggests accumulation across attempts; eligibility depends on prediction-time definition
8	Reporting decision (campaign)	Phrase campaign as timing-dependent feature requiring clarification before modeling	Accept	Updated leakage paragraph to distinguish confirmed leakage (duration) vs timing-ambiguous feature (campaign)	Avoids overclaiming at EDA stage while documenting risk and the need for operational definition
#	**Step 3 Element**	**Agent Suggestion**	**Decision**	**Verification**	**Rationale**
1	Feature eligibility	Exclude duration (post-call variable)	Accept	Step 2 leakage EDA: strong outcome association; known after call	Prevent target leakage under pre-contact definition
2	Feature eligibility	Exclude campaign due to timing ambiguity (accumulates within campaign)	Accept	Step 2 leakage note + distribution shows accumulation; not guaranteed pre-first-contact	Avoid temporal ambiguity for “first-call” scoring
3	Final feature list check	Provide final 14 feature columns used in X	Accept	Compared printed feature list against dataset schema	Ensures only intended predictors remain after exclusion
4	Timing validity audit	Flag day/month/contact as “usable only if planned pre-contact”	Modify	Documented assumption: planned schedule/channel known at scoring time (or exclude in stricter set)	Makes pre-contact definition explicit and defensible
5	Proxy leakage screening	Rank remaining features by mutual information; flag suspiciously predictive variables	Accept	Checked top-ranked features are mostly historical campaign signals (poutcome/pdays/previous)	Confirms no obvious hidden leakage like duration remains
6	Robustness plan	Create stricter feature set dropping day/month/contact and compare	Plan	Defined stricter feature set (Set B) to allow sensitivity analysis under alternative timing assumptions.	Demonstrates active governance and not blindly accepting features
7	Target encoding	Keep original string labels	Modify	Converted to 0/1 and verified positive class proportion	Ensures consistent metric interpretation
#	**Step 4 Element**	**Agent Suggestion**	**Decision**	**Verification (Experiment)**	**Rationale**
1	Hyperparameter ranges	Proposed small ranges for LR/RF/HGB	Accept	Evaluated with 5-fold Stratified CV using PR-AUC	Compare models under consistent protocol
2	Search strategy	Full grid for RF/HGB	Modify	Switched to RandomizedSearchCV (fixed n_iter) and compared CV PR-AUC	Controls runtime while still evidence-based
3	LR solver choice	Try liblinear/lbfgs	Modify	Fixed solver to avoid sparse/efficiency issues; validated via CV PR-AUC	Keep tuning focused on impactful params
4	Model shortlist	LR + RF + HGB	Accept	Final shortlist based on mean CV PR-AUC (not intuition)	Meets baseline + modern approach requirement
5	Target encoding consistency	Implicitly assumed binary handling	Modify	Explicitly encoded target as 0/1; confirmed y_train/y_test unique values = {0,1}; recomputed CV and test PR-AUC	Ensure positive class consistency across CV and test; avoid label ambiguity
6	Test-set discipline	Not explicitly enforced	Strengthen protocol	Confirmed CV performed only on training split; evaluated test set exactly once after model selection	Prevent test leakage and inflated performance estimates
7	Generalization gap check	Not explicitly discussed	Add evaluation step	Compared CV PR-AUC (0.792) vs Test PR-AUC (0.777); gap ≈ 0.015	Assess overfitting risk; small, acceptable gap indicates stable generalization
8	Variability assessment	Provided mean CV score only	Extend analysis	Reviewed std across folds (≈0.009); compared stability across models	Prefer models with strong performance AND low variance
9	Stopping rule for tuning	Suggested further search possible	Deliberately stop tuning	Performance improvements marginal (<0.01 between top models)	Avoid over-tuning; control complexity within coursework scope
10	Model interpretability tradeoff	Not emphasized	Retain baseline	Logistic Regression kept despite lower score	Maintain interpretable benchmark for fairness & comparison
11	Search budget control	Full grid possible	Constrain search	Used RandomizedSearchCV with fixed n_iter for RF/HGB	Balance computational efficiency with empirical rigor

12	Pipeline consistency	Recreated preprocessing per model	Unify via shared pipeline	Used consistent ColumnTransformer across models	Ensure fair comparison under identical preprocessing
#	**Step 5 Element**	**Agent Suggestion**	**Decision**	**Verification (Experiment)**	**Rationale**
1	Tuning strategy definition	Use **train CV only** for model selection; keep **test untouched**	Accept	Ran **Stratified CV PR-AUC** to pick best model; explicitly printed "test untouched" policy	Prevents optimistic bias; aligns with coursework requirement (evidence-based selection)
2	Thresholding rule	Choose decision threshold by **maximizing F1 on validation set** (not on test)	Accept	Selected threshold = **0.3368** using **validation**; then applied once to test to compute confusion matrix + F1	Avoids using test to tune threshold; separates "model selection" vs "operating point"
3	Robust test evaluation	Report **test PR-AUC + ROC-AUC**, plus thresholded metrics (precision/recall/F1)	Accept	On test (single use): **PR-AUC=0.7922**, **ROC-AUC=0.7915**, Precision/Recall/F1 reported from confusion matrix	PR-AUC is primary; ROC-AUC is complementary; threshold metrics explain trade-offs
4	Calibration check	Add **reliability curve** + **Brier score** to assess probability quality	Accept	Plotted reliability curve; computed **Brier score = 0.1828** on test	Model isn't only "ranking"; calibration matters if probabilities drive targeting decisions
5	Failure mode / error analysis	Inspect **top false positives / false negatives** with their predicted probabilities and key features	Accept	Extracted top 20 FP/FN; compared distributions (e.g., pdays/previous/balance patterns)	Turns metrics into actionable insight (where the model is confidently wrong)
6	Metric / label handling	Evaluation assumed correct binary encoding	Modify	Encountered pos_label=1 error; verified label type; enforced explicit 0/1 encoding; re-ran successfully	Ensures correct metric computation and positive-class definition
7	Test-set protocol formalization	No explicit test discipline enforced in code	Strengthen protocol	Defined rule: test evaluated exactly once after model selection; debugging confined to train/val	Prevents accidental test leakage during experimentation